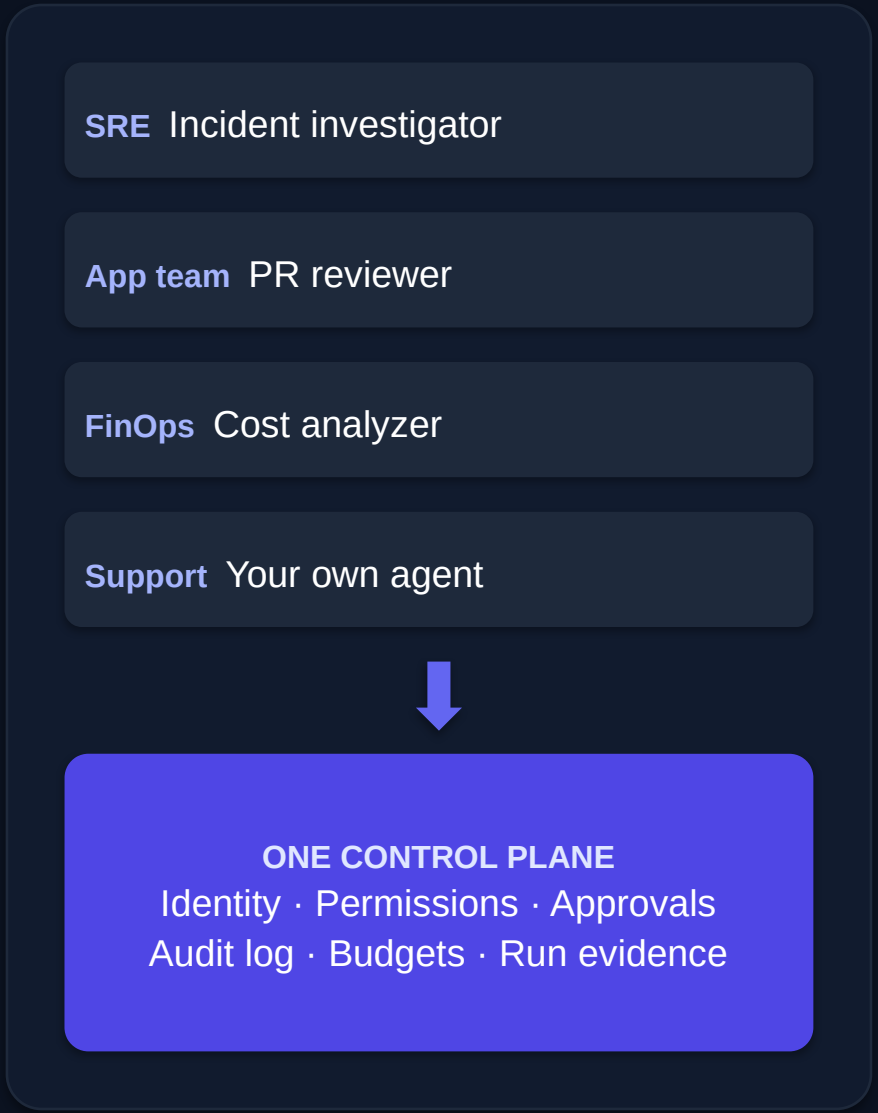
KOMODOR AGENTIC OPERATIONS PLATFORM

Your teams are already shipping agents.

Can you govern them?

Give every agent an identity, a boundary, an approval path and a record.



Echizeni Ikpah · Product Engineer

About me

Echizeni Ikpah

Product Engineer

With your team From your first agent to production.

With our product team Your feedback shapes what we build next.

WHY WE'RE HERE

Your teams launch AI agents on their own. Nobody sees them all or sets their limits.

WHAT KAOP GIVES YOUR PLATFORM TEAM

- One view of every agent
- Limits a system enforces
- Agents working as one workflow

Live today: a real incident and a real code review.

What to expect

About 30 minutes, then your questions.

5 min	**Your challenges**	Where agent adoption hurts today.
10 min	**How KAOP solves them**	Visibility, guardrails, workflows.
12 min	**Live: one real incident**	From alert to reviewed fix.
3 min	**A first step**	Shaped around your team.
15 min	**Your questions**	Anything, including the hard ones.

YOU WILL SEE, LIVE

1. **A real alert fires in Grafana**
   KAOP starts work on its own
2. **Two agents investigate**
   Root cause, with the evidence
3. **A person picks the fix**
   Agents propose. People decide.
4. **An agent reviews the fix**
   Risky changes caught before merge
5. **One record of every step**
   Who did what, and what it cost

Agents are spreading faster than your operating model

Each team's agent makes sense on its own. Nobody manages them together.

WHO OWNS IT?

Ownership

Agents appear team by team. Nobody has the full list.

WHAT CAN IT TOUCH?

Access

Each agent brings its own keys. Its limits live in a prompt.

WHAT DID IT DO?

Auditability

Alert, investigation and fix sit in three different tools.

40%+ **of agentic AI projects will be canceled by the end of 2027.**

Escalating costs, unclear business value or inadequate risk controls. Gartner, Inc., cited in Komodor's KAOP launch.

QUESTION Ownership, access or auditability: which hurts your team most today?

KAOP runs agents like production services

Each problem you have today, and what KAOP does about it.

TODAY
Every team builds its agents from scratch.

BUILD

Start from 50+ ready agents

Or bring your own. Same controls.

TODAY
Alert, investigation and fix sit in three tools.

RUN

One workflow, one record

Alerts start the work. Agents hand off.

TODAY
An agent's limits live in its prompt.

GOVERN

Limits the platform enforces

Identity, permissions, approvals and audit.

TODAY
Nobody knows what agents cost or are worth.

OPTIMIZE

Cost per run, budgets that block

Keep what earns its cost.

KOMODOR'S AI SRE IN PRODUCTION WITH NEBIUS DELL CISCO PRUDENTIAL

A prompt is not a permission model

KAOP governs an agent the way it governs a person.

LAYER		WHAT YOU GET
1	**Identity**	Every action has a name on it
2	**Permissions**	Only the access the job needs
3	**Standards**	No agent ships without an owner
4	**Guardrails**	Risky calls blocked, every run
5	**Approvals**	A person decides the irreversible
6	**Audit log**	Every step on the record

IN TODAY'S LAB

Investigators
Read only. Enforced outside the agent.

Remediation
Proposes fixes. In this lab, a person applies them.

Code reviewer
Comments only. Can't approve or merge.

PRINCIPLE Prompts drift. Enforcement doesn't.

One record from the alert to the reviewed fix

What the platform team sees, without asking anyone.

Fleet

Every agent, its owner, its status.

Runs & evidence

Inputs, tool calls, reasoning, output, cost.

Waiting on you

Decisions that need a person.

Audit log

Who changed what, and when.

Cost

Cost per agent, with budgets.

Access explorer

What can this agent actually do?

RESULT Inventory, access and audit become a screen, not a spreadsheet.

One incident, two governed workflows

Built from KAOP catalog agents. No custom code.

WORKFLOW 1 · INCIDENT · owned by SRE · starts on a Grafana alert



WORKFLOW 2 · CODE REVIEW · owned by platform engineering · starts on any pull request



NOTICE Agents investigate, propose and review. People approve and change production.

Who can do what, and what enforces it

Least privilege per agent, enforced outside the agent.

Agent	Can read	Can change	Enforced by
Kubernetes RCA	Pods, events, logs, workloads	Nothing	Kubernetes RBAC on its own identity
Grafana investigator	Metrics, logs, alerts	Nothing but links to its evidence	KAOP gateway allow list: 28 read tools
Orchestrator	Findings	Nothing	May call only allowlisted agents
Remediation	The investigation	Nothing	Plan-only by design
Code reviewer	The PR diff	One comment	GitHub app scopes

VERIFIED IN THIS LAB

Kubernetes RCA, checked with kubectl: reads pods, events, logs. Denied secrets, exec, delete, patch, scale.
Grafana investigator: any tool outside its allow list comes back "Tool not found".

Live: one bad change, one Redis outage

Real cluster, real alert, real agents.

BREAK

1
Bad change
Liveness checks Redis.

2
Redis goes down
Started when we began.

3
Amplified
Kubernetes kills the healthy API.

INVESTIGATE

4
Detect
Grafana alerts KAOP.

5
Investigate
Root cause, confidence, evidence.

DECIDE AND FIX

6
Decide
I approve, then apply.

7
Prevent
One-line PR, reviewed by an agent.

MEASURED ON THIS LAB

Bad check API killed 5 times in 5 minutes, alert fired

Fixed check 0 kills, alert stayed quiet

WATCH FOR Redis is the trigger. The health check is the root cause.

Autonomy, earned one step at a time

From human in the loop to human on the loop. Each step is a permission you grant.

1 TODAY

Investigate

Root cause and evidence, minutes after the alert.

2 TODAY

Propose

The fix, with undo and a check.

3 NEXT

Approve to execute

A person approves. An agent applies.

4 LATER

Pre-approved actions

Low-risk fixes run on their own. People review.

GOAL Lower time to resolution, fewer 3 a.m. pages, engineers back on the roadmap.

Let's shape the first step together

Built around your team, at your pace.

ALIGN

A scoping session

Agree the workflow, the owner, how we measure success, and the guardrails.

PROVE

One workflow

One cluster. Read-only. Every change approved by a person.

EXPAND

At your pace

More workflows, teams and autonomy, once the results hold.

FROM YOUR TEAM

- A driver who owns it day to day
- The budget owner, who decides when to expand
- Two or three dates for the scoping session

FROM KOMODOR

- A product engineer, hands-on with setup
- Catalog agents, ready to run
- A joint review of the results

NEXT Pick a date for the scoping session with your driver and budget owner.

Thank you

What you saw today

1. **One view of every agent**
   Who owns it, what it did, what it cost.

2. **Limits a system enforces**
   Kubernetes and KAOP said no. Not a prompt.

3. **Agents working as one workflow**
   Alert to reviewed fix, on one record.

NEXT STEP

A scoping session with your driver and budget owner.

Echizeni Ikpah
Product Engineer

Questions?

Appendix

Reference material

A Sources for this presentation

B Common questions

C Measured results

D Cost and guardrails, tested

Sources for this presentation

Every claim traces to Komodor's material or to this lab.

Komodor website and KAOP launch	Positioning: build, run, govern, optimize. The Gartner figure.
KAOP docs: security and governance	The six layers, approvals, audit log, Access explorer.
KAOP docs: agent catalog	What each agent reads, changes, and where it stops.
KAOP docs: incident module	The alert-to-remediation workflow.
This lab	A real fault, a real alert, permissions verified, fix measured.

APPROACH Start from your challenges, prove it on a real incident, then agree the first step together.

Common questions

“We'll build this ourselves.”	You can build an agent. The hard part is identity, approvals, audit, budgets and evaluation around every agent, for every team. That's the platform.
“Komodor already has Klaudia. Why KAOP?”	Klaudia is Komodor's AI SRE. KAOP is the control center for all your agents: Klaudia-based, Komodor's and your own.
“What if the agent is wrong?”	It separates what it saw from what it guessed, every change needs a person, and KAOP can grade runs with an evaluator.
“How long does onboarding take?”	Catalog agents deploy in minutes. The real time goes into agreeing the workflow, the owner and the guardrails, which is why I'd start with a scoping session.
“What does it cost?”	Every run shows its cost, and budgets block new runs at a limit. We'd agree a budget in the scoping session and track cost per useful investigation.

Measured results

The same Redis outage, measured both ways.

	Liveness checks Redis (bad)	Liveness checks the process (fix)
Times Kubernetes killed the healthy API	5 in a 5-minute outage	0
Liveness failures	Yes	None
Grafana restart alert	Fired at +3 minutes	Stayed Normal
API back after Redis returned	4.8 s to 3.6 min (restart back-off)	6.5 s

Measured on 2026-10-24 and 2026-10-25 on a single-node kind cluster. With the bad check, recovery waits on Kubernetes' restart back-off (10 s, doubling up to 5 minutes), so it depends on how long the outage lasted.

Cost and guardrails, tested

Each change made on this lab, then tested with real runs.

Change	Test	Result
Grafana agent: only the 28 read tools it needs	Same alert, 5 runs with and 5 without	About $0.68 → $0.53 a run. Same root cause
One standards doc for every incident agent	Same alert, 5 runs without the doc	Cause found in 10 of 10 runs with it, 1 of 5 without
Workflow budget: $40 a day, blocks new runs	Cap lowered to $0.01, two alerts sent	Second alert recorded but not investigated, with the reason
Grafana agent: allow list instead of a deny list	Asked for tools outside the list	"Not permitted" / "Tool not found"
Owner and team label on every agent	Fleet filtered by owner	Every agent has an owner to call

Measured on 2026-10-25 and 2026-10-26 (UTC). Identical runs vary by about 20 cents, so read the cost result as a direction, not a benchmark.